

# Multiplier Design Comparative Analysis (Pipelined v. Combinational)

Summer 2025

Designing a systolic array and TinyNPU taught me the importance of multiply-accumulate (MAC) processing elements in machine learning hardware. Thus, I designed a pipelined multiplier for signed INT8 using the Baugh-Wooley algorithm and compared it to a combinational (\*) multiplier in terms of power, performance, area, and other metrics.

## Background

### Baugh-Wooley Multiplication

The Baugh-Wooley method is useful for multiplying signed (i.e., two's complement) integers. Below is an example of multiplying two signed 4-bit integers.

$$\begin{array}{ccccccccc} & 0 & 0 & 0 & | & \overline{a_3 b_0} & a_2 b_0 & a_1 b_0 & a_0 b_0 & c_{p_0} \\ + & 0 & 0 & 0 & \overline{a_3 b_1} & a_2 b_1 & a_1 b_1 & a_0 b_1 & & c_{p_1} \\ + & 0 & 0 & \overline{a_3 b_2} & \overline{a_2 b_2} & a_1 b_2 & a_0 b_2 & & & c_{p_2} \\ + & | & \overline{a_3 b_3} & \overline{a_2 b_3} & \overline{a_1 b_3} & \overline{a_0 b_3} & & & & c_{p_3} \\ \hline & z_7 & z_6 & z_5 & z_4 & z_3 & z_2 & z_1 & z_0 & \end{array}$$

Essentially, the Baugh-Wooley method allows multiplying two signed integers without requiring excess hardware for sign-extended partial products. This is made possible through [several steps of boolean algebra](#) that simplifies to (1) inverting the MSB of each line of partial products, (2) adding constants to the first and last line of partial products while inverting the latter, and (3) accumulating all the partial products.

### INT8

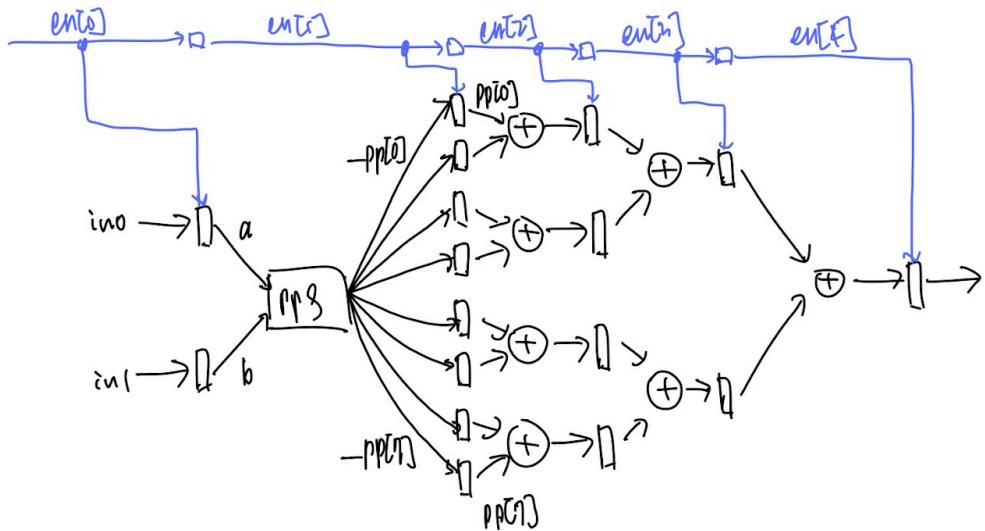
INT8 stands for 8-bit signed integers that are often used for quantizing deep learning model parameters. INT8 is often used for inference computations because it is both time- and space-efficient than large bitwidths like 32-bit floating point numbers (FP32). Thus, it is critical to have efficient hardware designs for multiplying INT8 numbers involved in MAC operations.

Note that multiplying two INT8 numbers yield an INT16 number. When performing intermediate MAC operations, 32-bit accumulators are often used to add those INT16 numbers to preserve precision until saturating to INT8 as needed.

# RTL

## Pipelined Multiplier

The pipelined multiplier for signed INT8 has five stages in the following order: input stage, partial product generation (PPG) stage, and three accumulation reduction tree stages. Below is a diagram of the pipelined multiplier.



Note that all logic between each pipeline stage is combinational (i.e., no stalling). Also, the pipelined multiplier has an enable signal going through the pipeline along with the two inputs. This assumes that the pipelined multiplier is a building block included in a larger system with an overarching control logic for correctly enabling the multiplier as needed.

## Combinational Multiplier

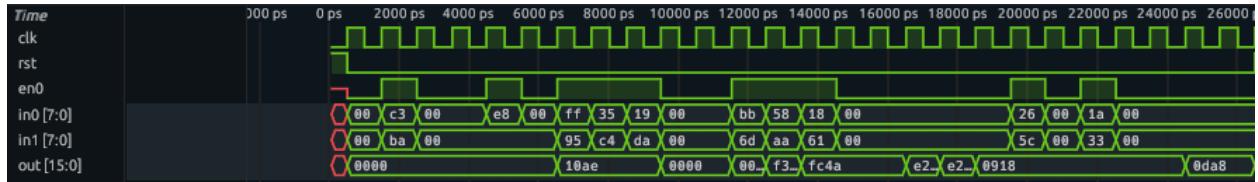
The combinational multiplier simply uses the  $*$  operator triggered by a rising clock edge.

## Design Verification

The **pipelined multiplier** was simulated using Cocotb (Synopsys VCS) and the following test cases. Note that various number distributions (e.g., full domain, full positive domain, full negative domain, small/large) positive/negative domain) were used 10,000 times each for generating INT8 input pairs.

- Simple directed multiplications (e.g., alternating  $0 \times 0 = 0$  and  $1 \times 1 = 1$  transactions)
- Directed input stream with constant enable and random INT8 input pairs
- Random input stream with random enable and random INT8 input pairs

The waveform shows a random input stream test case.



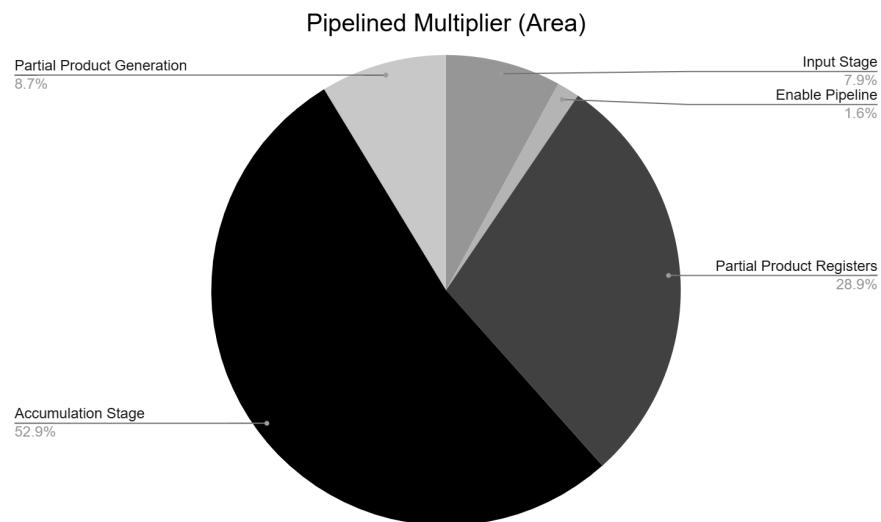
## Physical Design

*Technical details beyond what is provided below cannot be shared under the non-disclosure agreement between Cornell Custom Silicon Systems, MPW (Multi-Project-Wafer) University Service (MUSE), and Taiwan Semiconductor Manufacturing Company (TSMC).*

The MUSE TSMC 180 nm process node was used to collect power, performance, and area metrics for the pipelined and combinational multipliers.

### Area

The **pipelined multiplier** requires approximately  $16,590 \mu\text{m}^2$  of space where much of this area is allocated for the accumulation reduction tree and the partial product generator and registers.



The **combinational multiplier** requires approximately  $6,900 \mu\text{m}^2$  of space.

### Performance

The **pipelined multiplier** can operate at a clock period of 3 ns (333 MHz) where the critical path goes through the combinational adder of the last accumulation stage with a slack of 0.394 ns.

Path 1: MET Setup Check with Pin g\_s2\_reg\_0\_acc\_reg/q\_reg\_15\_/CP  
Endpoint: g\_s2\_reg\_0\_acc\_reg/q\_reg\_15\_/D (v) checked with leading edge of 'Multiplier\_INT8\_clk'

```

Beginpoint: g_s1_reg_1__acc_reg/q_reg_4_/Q (v) triggered by leading edge of
'Multiplier_INT8_clk'
Path Groups: {Multiplier_INT8_clk}
Analysis View: analysis_typical
Other End Arrival Time 0.000
- Setup 0.054
+ Phase Shift 3.000
+ CPPR Adjustment 0.000
= Required Time 2.946
- Arrival Time 2.551
= Slack Time 0.394
Clock Rise Edge 0.000
+ Clock Network Latency (Ideal) 0.000
= Beginpoint Arrival Time 0.000
+-----+
| Instance | Arc | Cell | Delay | Arrival | Required |
| | | | | Time | |
+-----+
| g_s1_reg_1__acc_reg/q_reg_4_| CP ^ | | 0.000 | 0.394 | |
| g_s1_reg_1__acc_reg/q_reg_4_| CP ^ -> Q v | xxxxxxxxxx | 0.270 | 0.270 | 0.664 |
| g_s2_adder_0__acc_adder/U5 | A2 v -> Z v | xxxxxxxxxx | 0.120 | 0.390 | 0.784 |
| g_s2_adder_0__acc_adder/intadd_1_U9 | CI v -> CO v | xxxxxxxxxx | 0.192 | 0.582 | 0.976 |
| g_s2_adder_0__acc_adder/intadd_1_U8 | CI v -> CO v | xxxxxxxxxx | 0.200 | 0.782 | 1.176 |
| g_s2_adder_0__acc_adder/intadd_1_U7 | CI v -> CO v | xxxxxxxxxx | 0.211 | 0.993 | 1.388 |
| g_s2_adder_0__acc_adder/intadd_1_U6 | CI v -> CO v | xxxxxxxxxx | 0.198 | 1.192 | 1.586 |
| g_s2_adder_0__acc_adder/intadd_1_U5 | CI v -> CO v | xxxxxxxxxx | 0.193 | 1.385 | 1.779 |
| g_s2_adder_0__acc_adder/intadd_1_U4 | CI v -> CO v | xxxxxxxxxx | 0.192 | 1.577 | 1.971 |
| g_s2_adder_0__acc_adder/intadd_1_U3 | CI v -> CO v | xxxxxxxxxx | 0.213 | 1.790 | 2.184 |
| g_s2_adder_0__acc_adder/intadd_1_U2 | CI v -> CO v | xxxxxxxxxx | 0.175 | 1.964 | 2.359 |
| g_s2_adder_0__acc_adder/U2 | A1 v -> Z v | xxxxxxxxxx | 0.117 | 2.081 | 2.476 |
| g_s2_adder_0__acc_adder/U6 | A2 v -> ZN ^ | xxxxxxxxxx | 0.063 | 2.144 | 2.539 |
| g_s2_adder_0__acc_adder/U7 | A2 ^ -> ZN v | xxxxxxxxxx | 0.186 | 2.330 | 2.724 |
| g_s2_reg_0__acc_reg/U19 | B2 v -> Z v | xxxxxxxxxx | 0.221 | 2.551 | 2.945 |
| g_s2_reg_0__acc_reg/q_reg_15_| D v | xxxxxxxxxx | 0.000 | 2.551 | 2.946 |
+-----+

```

**The combinational multiplier can operate at a clock period of 4 ns (250 MHz) with a slack of 0.214 ns.**

```

Path 1: MET Setup Check with Pin out_reg_14_/CP
Endpoint: out_reg_14_/D (^) checked with leading edge of 'StarMultiplier_clk'
Beginpoint: in0[3] (v) triggered by leading edge of 'StarMultiplier_clk'
Path Groups: {StarMultiplier_clk}
Analysis View: analysis_typical
Other End Arrival Time 0.000
- Setup 0.062
+ Phase Shift 4.000
+ CPPR Adjustment 0.000
= Required Time 3.938
- Arrival Time 3.724
= Slack Time 0.214
Clock Rise Edge 0.000
+ Input Delay 0.005
= Beginpoint Arrival Time 0.005
+-----+
| Instance | Arc | Cell | Delay | Arrival | Required |
| | | | | Time | |
+-----+
| | in0[3] v | | 0.005 | 0.219 | |
| FE_DBTC1_in0_3 | I v -> ZN ^ | xxxxxxxxxx | 0.038 | 0.044 | 0.257 |
| U60 | A1 ^ -> Z ^ | xxxxxxxxxx | 0.104 | 0.148 | 0.361 |
| U42 | I ^ -> ZN v | xxxxxxxxxx | 0.036 | 0.184 | 0.397 |
| U199 | A2 v -> ZN ^ | xxxxxxxxxx | 0.126 | 0.309 | 0.523 |
| U116 | B2 ^ -> Z ^ | xxxxxxxxxx | 0.143 | 0.452 | 0.666 |
| U175 | I ^ -> ZN v | xxxxxxxxxx | 0.040 | 0.492 | 0.706 |
| U65 | A1 v -> Z v | xxxxxxxxxx | 0.121 | 0.613 | 0.827 |
| U141 | I v -> ZN ^ | xxxxxxxxxx | 0.041 | 0.655 | 0.868 |
| U52 | B ^ -> Z ^ | xxxxxxxxxx | 0.095 | 0.750 | 0.963 |
| U138 | I ^ -> ZN v | xxxxxxxxxx | 0.050 | 0.800 | 1.014 |
+-----+

```

U56	A v -> ZN ^	xxxxxxxxxxxx	0.213	1.013	1.227
U307	A ^ -> ZN v	xxxxxxxxxxxx	0.230	1.243	1.457
U68	I v -> ZN ^	xxxxxxxxxxxx	0.039	1.282	1.496
U133	A2 ^ -> ZN v	xxxxxxxxxxxx	0.052	1.334	1.548
U137	A2 v -> Z v	xxxxxxxxxxxx	0.108	1.442	1.656
U292	A2 v -> ZN ^	xxxxxxxxxxxx	0.076	1.518	1.732
U296	C ^ -> ZN v	xxxxxxxxxxxx	0.193	1.711	1.925
U257	A v -> ZN ^	xxxxxxxxxxxx	0.253	1.964	2.178
intadd_1_U2	A ^ -> CO ^	xxxxxxxxxxxx	0.271	2.234	2.448
U55	B ^ -> ZN v	xxxxxxxxxxxx	0.093	2.327	2.541
intadd_2_U2	CI v -> CO v	xxxxxxxxxxxx	0.174	2.502	2.715
U253	B v -> ZN ^	xxxxxxxxxxxx	0.278	2.780	2.993
intadd_0_U2	A ^ -> CO ^	xxxxxxxxxxxx	0.268	3.048	3.262
U153	I ^ -> ZN v	xxxxxxxxxxxx	0.034	3.082	3.296
U300	C v -> ZN ^	xxxxxxxxxxxx	0.281	3.364	3.577
U66	A3 ^ -> Z v	xxxxxxxxxxxx	0.220	3.583	3.797
U54	A2 v -> ZN ^	xxxxxxxxxxxx	0.141	3.724	3.938
out_reg_14_	D ^	xxxxxxxxxxxx	0.000	3.724	3.938

## Power

The **pipelined multiplier** consumes 10.3 mW of power.

Total Power

Total Internal Power:	9.00352961	87.4074%
Total Switching Power:	1.29705423	12.5920%
Total Leakage Power:	0.00006185	0.0006%
Total Power:	10.30064570	

Group	Internal Power	Switching Power	Leakage Power	Total Power	Percentage (%)
Sequential	7.141	0.4968	2.811e-05	7.638	74.15
Macro	0	0	0	0	0
IO	0	0	0	0	0
Combinational	1.863	0.8002	3.375e-05	2.663	25.85
Clock (Combinational)	0	0	0	0	0
Clock (Sequential)	0	0	0	0	0
<b>Total</b>	<b>9.004</b>	<b>1.297</b>	<b>6.185e-05</b>	<b>10.3</b>	<b>100</b>

Rail	Voltage	Internal Power	Switching Power	Leakage Power	Total Power	Percentage (%)
VDD	1.8	9.004	1.297	6.185e-05	10.3	100

Hierarchy	Internal Power	Switching Power	Leakage Power	Total Power	Percentage (%)
g_en_reg_1_en_reg	0.0475	0.02737	1.991e-07	0.07487	0.7269
g_en_reg_2_en_reg	0.04518	0.01815	2.503e-07	0.06333	0.6148
g_en_reg_3_en_reg	0.04273	0.001992	2.459e-07	0.04472	0.4342
g_en_reg_4_en_reg	0.04049	0.0006269	2.437e-07	0.04111	0.3991
a_reg	0.4045	0.1057	2.449e-06	0.5102	4.953
b_reg	0.4108	0.08573	2.506e-06	0.4965	4.82
ppg	0.1694	0.2149	4.662e-06	0.3843	3.731
g_pp_reg_0_pp_reg	0.4228	0.0433	2.346e-06	0.4661	4.525
g_pp_reg_1_pp_reg	0.3629	0.02907	1.998e-06	0.392	3.805
g_pp_reg_2_pp_reg	0.3624	0.03315	1.98e-06	0.3956	3.84

g_pp_reg_3_pp_reg	0.3727	0.03879	2.07e-06	0.4115	3.994
g_pp_reg_4_pp_reg	0.3629	0.0295	1.998e-06	0.3924	3.809
g_pp_reg_5_pp_reg	0.3727	0.04086	2.07e-06	0.4135	4.015
g_pp_reg_6_pp_reg	0.3723	0.0361	2.069e-06	0.4084	3.965
g_pp_reg_7_pp_reg	0.4566	0.05111	2.582e-06	0.5077	4.929
g_s0_adder_0_acc_adder	0.08522	0.009178	1.691e-06	0.0944	0.9165
g_s0_adder_2_acc_adder	0.08206	0.0114	1.711e-06	0.09346	0.9073
g_s0_adder_4_acc_adder	0.08165	0.009244	1.711e-06	0.0909	0.8824
g_s0_adder_6_acc_adder	0.1201	0.02798	1.678e-06	0.1481	1.438
g_s0_reg_0_acc_reg	0.4766	0.05031	2.643e-06	0.5269	5.115
g_s0_reg_1_acc_reg	0.4714	0.04964	2.389e-06	0.5211	5.058
g_s0_reg_2_acc_reg	0.4772	0.04935	2.662e-06	0.5265	5.111
g_s0_reg_3_acc_reg	0.4963	0.05446	2.631e-06	0.5507	5.346
g_s1_adder_0_acc_adder	0.1375	0.01878	2.115e-06	0.1563	1.517
g_s1_adder_2_acc_adder	0.1419	0.02139	1.883e-06	0.1632	1.585
g_s1_reg_0_acc_reg	0.6343	0.07393	3.387e-06	0.7082	6.876
g_s1_reg_1_acc_reg	0.595	0.07625	3.167e-06	0.6713	6.517
g_s2_adder_0_acc_adder	0.1688	0.024	2.365e-06	0.1928	1.872
g_s2_reg_0_acc_reg	0.7898	0.06484	4.149e-06	0.8547	8.297

The **combinational multiplier** consumes 2.5 mW of power.

#### Total Power

---

Total Internal Power:	1.91007316	77.4983%
Total Switching Power:	0.55456115	22.5005%
Total Leakage Power:	0.00003071	0.0012%
Total Power:	2.46466502	

---

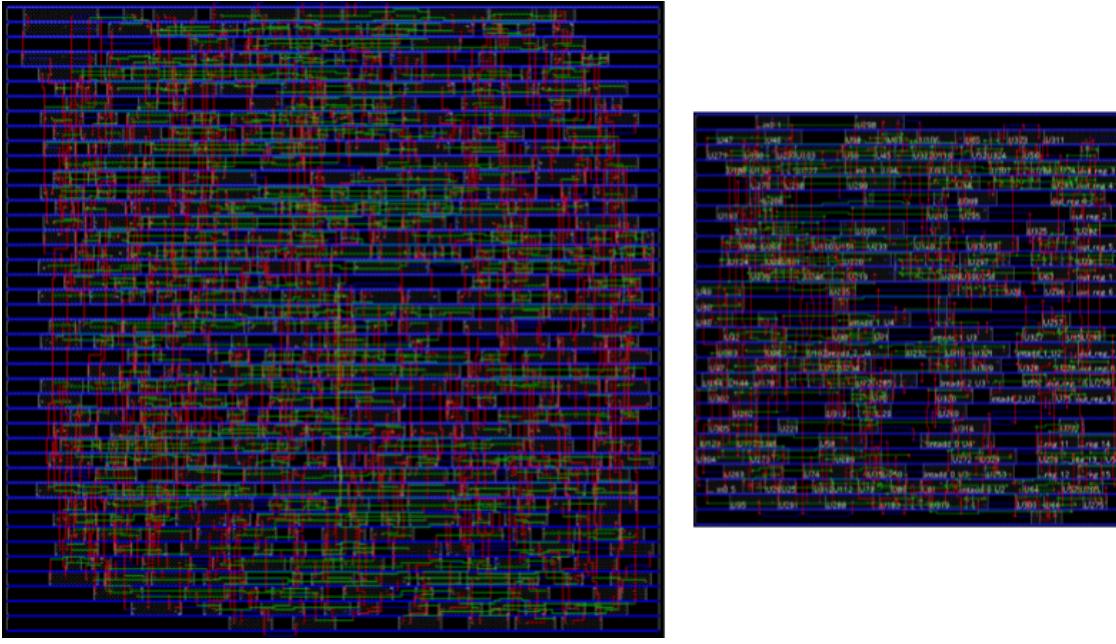
Group	Internal Power	Switching Power	Leakage Power	Total Power	Percentage (%)
Sequential	0.5631	0.0213	2.476e-06	0.5844	23.71
Macro	0	0	0	0	0
IO	0	0	0	0	0
Combinational	1.347	0.5333	2.823e-05	1.88	76.29
Clock (Combinational)	0	0	0	0	0
Clock (Sequential)	0	0	0	0	0
Total	1.91	0.5546	3.071e-05	2.465	100

---

Rail	Voltage	Internal Power	Switching Power	Leakage Power	Total Power	Percentage (%)
VDD	1.8	1.91	0.5546	3.071e-05	2.465	100

## Layout

Below is the final layout of the pipelined (left) and combinational (right) multipliers without filler cells (roughly in relative scale).



## Comparative Analysis

Metric	Pipelined Multiplier	Combinational (*) Multiplier
Process Node	<i>TSMC 180 nm</i>	
I/O	<i>INT8 x INT8 = INT16</i>	
Area	<i>16,590 μm<sup>2</sup></i>	<i>6,900 μm<sup>2</sup></i>
Power	<i>10.3 mW</i>	<i>2.5 mW</i>
Performance	<i>333 MHz (3 ns)</i>	<i>250 MHz (4 ns)</i>
Implementation	<i>High complexity</i>	<i>Low complexity</i>
Physical Design	<i>Dependent on RTL</i>	<i>Dependent on synthesis tool</i>

The combinational multiplier consumes 146% and 312% less area and power, respectively, than the pipelined multiplier due to the numerous registers between each pipeline stage and repeated adders in the accumulation reduction tree. However, the pipelined multiplier's critical path only goes through a combinational adder, which intuitively expends less logic and time than the combinational multiplier to finish its compute per cycle. As a result, the pipelined multiplier's clock frequency can be 33% higher than that of the combinational multiplier (note that the \* is triggered by a rising clock edge).

The pipelined multiplier has higher implementation complexity than the combinational multiplier because the former requires various boolean and RTL considerations while the latter simply uses the \* operator. This means that the pipelined multiplier's final physical layout is highly

dependent on the architect's choices while the combinational multiplier is entirely dependent on the synthesis tool. Thus, designing a pipelined multiplier may provide more control and knowledge of the hardware (at the expense of increased production time) than designing a combinational multiplier.

## Summary

- The combinational (\*) multiplier is best for applications requiring area- and power-efficient single-cycle operations at the cost of lower throughput and losing control over physical synthesis.
- The pipelined multiplier is best for applications requiring high throughput and more control over physical synthesis at the cost of additional area, power, and implementation complexity.